

# SQL – Data Manipulation Language

---

Carla Teixeira Lopes

Bases de Dados

Mestrado Integrado em Engenharia Informática e Computação, FEUP

Based on Jennifer Widom slides

# Agenda

---

Introduction

The JOIN family of operators

Basic SQL Statement

Aggregation

Table Variables and Set Operators

Null values

Subqueries in WHERE clauses

Data Modification statements

Subqueries in FROM and SELECT clauses

# SQL

---

Stands for **S**tructured **Q**uery **L**anguage

Pronounced “sequel”

Supported by all major commercial database systems

Standardized – many features over time

Interactive via GUI or prompt, or embedded in programs

Declarative, based on relational algebra

# SQL History

---

**1970** “A Relational Model of Data for Large Shared Data Banks” by Edgar Codd

Early **70's** SEQUEL Developed at IBM by Donald Chamberlin e Raymond Boyce

**1979** First commercial version by Relational Software (now Oracle)

**1986** SQL-86 and SQL-87. Ratified by ANSI and ISO

**1989** SQL-891992 SQL-92. Also known as SQL2

**1999** SQL:1999. Also known as SQL3. Includes regular expressions, recursive queries, triggers, non-scalar data types and some object-oriented expressions

**2003** SQL:2003 XML support and auto-generated values

**2006** SQL:2006 XQuery support

**2008** SQL:2008

**2011** SQL:2011

# SQL is a ...

---

## Data Definition Language (DDL)

- Define relational schemata

- Create/alter/delete tables and their attributes

## Data Manipulation Language (DML)

- Insert/delete/modify tuples in tables

- Query one or more tables

# Standard

---

Many standards out there

Database management systems implement something similar, but not identical to the standard for SQL

These slides will try to adhere to the standard as much as possible

Primarily the SQL2 standard and some constructs from the SQL3 standard

Sometimes we'll talk specifically about SQL as understood by SQLite

# Agenda

---

Introduction

The JOIN family of operators

Basic SELECT Statement

Aggregation

Table Variables and Set Operators

Null values

Subqueries in WHERE clauses

Data Modification statements

Subqueries in FROM and SELECT clauses

# The Basic SELECT Statement (SFW)

---

SELECT  $A_1, A_2, \dots, A_n$   $\longrightarrow$  What to return

FROM  $R_1, R_2, \dots, R_m$   $\longrightarrow$  Identifies the relations to query

WHERE condition  $\longrightarrow$  Combines and filters relations

$\rightarrow$  Produto cartesiano

What's the equivalent in Relational Algebra?

$\pi_{A_1, \dots, A_n} (\sigma_{condition} (R_1 \times \dots \times R_m))$

The result is a relation with the schema  $A_1, A_2, \dots, A_n$



# SQL is compositional

---

Relational query languages are compositional

When a query is run over relations, the result is a relation

The schema of the obtained relation is the set of attributes that are returned

The output of one query can be used as the input to another (nesting)

**This is extremely powerful**

# College Admission Database

Apply

<u>sID</u>	<u>cName</u>	<u>major</u>	<u>dec</u>
123	Stanford	CS	Y
123	Stanford	EE	N
123	Berkeley	CS	Y
123	Cornell	EE	Y
234	Berkeley	biology	N
345	MIT	bioengineering	Y
345	Cornell	bioengineering	N
345	Cornell	CS	Y
345	Cornell	EE	N
678	Stanford	history	Y
987	Stanford	CS	Y
987	Berkeley	CS	Y
876	Stanford	CS	Y
876	MIT	biology	Y
876	MIT	marine biology	N
765	Stanford	history	Y
765	Cornell	history	N
765	Cornell	psychology	Y
543	MIT	CS	N

College

<u>cName</u>	<u>state</u>	<u>enr</u>
Stanford	CA	15000
Berkeley	CA	36000
MIT	MA	10000
Cornell	NY	21000

Student

<u>sID</u>	<u>sName</u>	<u>GPA</u>	<u>HS</u>
123	Amy	3.9	1000
234	Bob	3.6	1500
345	Craig	3.5	500
456	Doris	3.9	1000
567	Edward	2.9	2000
678	Fay	3.8	200
789	Gary	3.4	800
987	Helen	3.7	800
876	Irene	3.9	400
765	Jay	2.9	1500
654	Amy	3.9	1000
543	Craig	3.4	2000

# SQL query with one relation

---

```
SELECT sID, sName, GPA
```

```
FROM Student
```

```
WHERE GPA > 3.6;
```

Interrogação com uma única relação

este atributo  
não tem de estar  
no resultado

sID	sName	GPA
123	Amy	3.9
456	Doris	3.9
678	Fay	3.8
987	Helen	3.7
876	Irene	3.9
654	Amy	3.9

Not necessary to include GPA in the result even if we filter on the GPA

College(cName, state, enr)

Student(sID, sName, GPA, sizeHS)

Apply(sID, cName, major, decision)

# A Few Details

---

SQL commands are case insensitive → 0 mesmo para os nomes das tabelas

Same: SELECT, Select, select

Same: Student, student

Values are not → Cuidado

Different: 'Stanford', 'stanford'

→ Use single quotes for constants

'abc' - yes

"abc" - no

# SQL query combining two relations

```
SELECT sName, major
FROM Student, Apply
WHERE Student.sID=Apply.sID;
```

This would happen automatically in a natural join of RA

What does it computes?

estudantes que se candidataram

Devido ao SQL ser baseado em múltiplos sets (Bugs)

Duplicate values

College(cName, state, enr)
Student(sID, sName, GPA, sizeHS)
Apply(sID, cName, major, decision)

sName	major
Amy	CS
Amy	EE
Amy	CS
Amy	EE
Bob	biology
Craig	bioengineering
Craig	bioengineering
Craig	CS
Craig	EE
Fay	history
Helen	CS
Helen	CS
Irene	CS
Irene	biology
Irene	marine biology
Jay	history
Jay	history
Jay	psychology
Craig	CS

# SQL query excluding duplicate values

Para remover conjuntos (sName, major) duplicados  
↪

```
SELECT DISTINCT sName, major  
FROM Student, Apply  
WHERE Student.sID=Apply.sID;
```

sName	major
Amy	CS
Amy	EE
Bob	biology
Craig	bioengineering
Craig	CS
Craig	EE
Fay	history
Helen	CS
Irene	CS
Irene	biology
Irene	marine biology
Jay	history
Jay	psychology

No duplicate values



College( <u>cName</u> , state, enr)
Student( <u>sID</u> , sName, GPA, sizeHS)
Apply( <u>sID</u> , <u>cName</u> , <u>major</u> , decision)

## Another SQL query combining two relations


---

SELECT sName, GPA, decision


FROM Student, Apply

WHERE Student.sID=Apply.sID AND sizeHS<1000 AND  
major='CS' AND cName='Stanford';

*Podemos juntar  
con diques*



What does it compute?



sName	GPA	decision
Helen	3.7	Y
Irene	3.9	N

College(cName, state, enr)

Student(sID, sName, GPA, sizeHS)

Apply(sID, cName, major, decision)

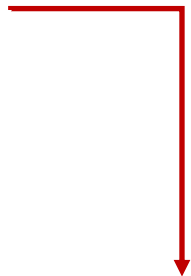
# One more SQL query combining two relations

---

SELECT cName  
FROM College, Apply  
WHERE College.cName=Apply.cName AND  
enr>20000 AND major='CS';

*→ Apenas 1 atributo → Vai dar ERRO, pois não especificamos de que tabela é que vem o cName, tal como aqui \**

*→ Duas tabelas*



SQLite Error: ambiguous  
column name: cName

What does it compute?

How can we correct it?

College(cName, state, enr)

Student(sID, sName, GPA, sizeHS)

Apply(sID, cName, major, decision)



# One more SQL query combining two relations

---

SELECT College.cName

FROM College, Apply

WHERE College.cName=Apply.cName AND  
enr>20000 AND major='CS';

utilizar DISTINCT }

cName
Berkeley
Berkeley
Cornell



How can we eliminate duplicates?

College(cName, state, enr)

Student(sID, sName, GPA, sizeHS)

Apply(sID, cName, major, decision)

# Order of the results

---

SQL is based on an unordered model

The order of the results may change each time we run a query

We can ask for a result to be sorted, in ascending or descending order, by an attribute or set of attributes

{ ORDER BY <attributes> ASC  
ORDER BY <attributes> DESC

Ordering is ascending, unless you specify the DESC keyword

# Order of the results

College(cName, state, enr)

Student(sID, sName, GPA, sizeHS)

Apply(sID, cName, major, decision)

```
SELECT Student.sID, sName, GPA,  
       Apply.cName, enr  
FROM   Student, College, Apply  
WHERE  Student.sID=Apply.sID AND  
       Apply.cName=College.cName;
```

```
SELECT Student.sID, sName, GPA,  
       Apply.cName, enr  
FROM   Student, College, Apply  
WHERE  Student.sID=Apply.sID AND  
       Apply.cName=College.cName  
ORDER BY GPA DESC;
```

sID	sName	GPA	cName	enr
123	Amy	3.9	Stanford	15000
123	Amy	3.9	Stanford	15000
123	Amy	3.9	Berkeley	36000
123	Amy	3.9	Cornell	21000
234	Bob	3.6	Berkeley	36000
345	Craig	3.5	MIT	10000
345	Craig	3.5	Cornell	21000
345	Craig	3.5	Cornell	21000
345	Craig	3.5	Cornell	21000
678	Fay	3.8	Stanford	15000

fluda  
→  
a ordem

sID	sName	GPA	cName	enr
123	Amy	3.9	Stanford	15000
123	Amy	3.9	Stanford	15000
123	Amy	3.9	Berkeley	36000
123	Amy	3.9	Cornell	21000
876	Irene	3.9	Stanford	15000
876	Irene	3.9	MIT	10000
876	Irene	3.9	MIT	10000
678	Fay	3.8	Stanford	15000
987	Helen	3.7	Stanford	15000
987	Helen	3.7	Berkeley	36000

# Order of the results

---

SELECT Student.sID, sName, GPA, Apply.cName, enr

FROM Student, College, Apply

WHERE Student.sID=Apply.sID AND Apply.cName=College.cName

ORDER BY GPA DESC, enr;

*Primeiro ordena por GPA e depois por enr*

sID	sName	GPA	cName	enr
876	Irene	3.9	MIT	10000
876	Irene	3.9	MIT	10000
123	Amy	3.9	Stanford	15000
123	Amy	3.9	Stanford	15000
876	Irene	3.9	Stanford	15000
123	Amy	3.9	Cornell	21000
123	Amy	3.9	Berkeley	36000
678	Fay	3.8	Stanford	15000
987	Helen	3.7	Stanford	15000
987	Helen	3.7	Berkeley	36000



Descending GPA as primary sort order and, within each of those, ascending enrollment



College(cName, state, enr)  
Student(sID, sName, GPA, sizeHS)  
Apply(sID, cName, major, decision)

# Like operator: pattern matching on strings

---

Built-in operator that allows string matching on attribute values

```
SELECT sID, major
```

```
FROM Apply
```

```
WHERE major like '%bio%';
```

*operador para comparar strings*

*CURSO tem de ter bio*

Match any major containing bio

sID	major
234	biology
345	bioengineering
345	bioengineering
876	biology
876	marine biology

% = any sequence of characters

\_ = any single character

College(cName, state, enr)

Student(sID, sName, GPA, sizeHS)

Apply(sID, cName, major, decision)

# Selecting all attributes

---

SELECT \*  seleccionar todos os atributos

FROM Apply

WHERE major like '%bio%';

sID	cName	major	decision
234	Berkeley	biology	N
345	MIT	bioengineering	Y
345	Cornell	bioengineering	N
876	MIT	biology	Y
876	MIT	marine biology	N

College(cName, state, enr)

Student(sID, sName, GPA, sizeHS)

Apply(sID, cName, major, decision)

# Arithmetic within SQL clauses

```
SELECT sID, sName, GPA, HS, GPA*(HS/1000)
FROM Student;
```

podemos ainda colocar  
cálculos aritméticos

sID	sName	GPA	HS	GPA*(HS/1000)
123	Amy	3.9	1000	3.9
234	Bob	3.6	1500	5.4
345	Craig	3.5	500	1.75
456	Doris	3.9	1000	3.9
567	Edward	2.9	2000	5.8
678	Fay	3.8	200	0.76
789	Gary	3.4	800	2.72
987	Helen	3.7	800	2.96
876	Irene	3.9	400	1.56
765	Jay	2.9	1500	4.35
654	Amy	3.9	1000	3.9
543	Craig	3.4	2000	6.8

Boosts GPA if student is from a big high school and reduces it, if he is from a small one

→ Can we improve this result?

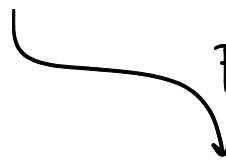
College(cName, state, enr)  
Student(sID, sName, GPA, sizeHS)  
Apply(sID, cName, major, decision)

# Renaming columns

---

```
SELECT sID, sName, GPA, HS, GPA*(HS/1000) AS scaledGPA
FROM Student;
```

permite renombrar los atributos



sID	sName	GPA	HS	scaledGPA
123	Amy	3.9	1000	3.9
234	Bob	3.6	1500	5.4
345	Craig	3.5	500	1.75
456	Doris	3.9	1000	3.9
567	Edward	2.9	2000	5.8
678	Fay	3.8	200	0.76
789	Gary	3.4	800	2.72
987	Helen	3.7	800	2.96
876	Irene	3.9	400	1.56
765	Jay	2.9	1500	4.35
654	Amy	3.9	1000	3.9
543	Craig	3.4	2000	6.8

College(cName, state, enr)  
Student(sID, sName, GPA, sizeHS)  
Apply(sID, cName, major, decision)



# Agenda

---

Introduction

The JOIN family of operators

~~Basic SQL Statement~~

Aggregation

Table Variables and Set Operators

Null values

Subqueries in WHERE clauses

Data Modification statements

Subqueries in FROM and SELECT clauses

# Table variables → Usadas nos froms para ter interrogações menos verbosas Para Renomear

---

Used in the FROM clause for two purposes

Make queries more readable

Rename relations when we have two instances of the same relation

```
SELECT Student.sID, sName, GPA, Apply.cName, enr
FROM   Student, College, Apply
WHERE  Student.sID=Apply.sID AND Apply.cName=College.cName;
```

```
SELECT S.sID, sName, GPA, A.cName, enr
FROM   Student S, College C, Apply A
WHERE  S.sID=A.sID AND A.cName=C.cName;
```

→ Not changing the  
result, only making  
the query more  
readable

# Table variables

---

How to list the pairs of students who have the same GPA?

Student(sID, sName, GPA, sizeHS)

```
SELECT S1.sID, S1.sName, S1.GPA, S2.sID, S2.sName, S2.GPA
FROM Student S1, Student S2
WHERE S1.GPA=S2.GPA;
```

*Tlesmo GPA*

How to list only pairs  
of different students?

*Diferentes IDs*

sID	sName	GPA	sID1	sName1	GPA1
123	Amy	3.9	123	Amy	3.9
123	Amy	3.9	456	Doris	3.9
123	Amy	3.9	876	Irene	3.9
123	Amy	3.9	654	Amy	3.9
234	Bob	3.6	234	Bob	3.6
345	Craig	3.5	345	Craig	3.5
456	Doris	3.9	123	Amy	3.9
456	Doris	3.9	456	Doris	3.9
...	...	...	...	...	...

# Table variables

---

```
SELECT S1.sID, S1.sName, S1.GPA, S2.sID, S2.sName, S2.GPA
FROM Student S1, Student S2
WHERE S1.GPA=S2.GPA AND S1.sID <> S2.sID;
```

sID	sName	GPA	sID1	sName1	GPA1
123	Amy	3.9	456	Doris	3.9
123	Amy	3.9	876	Irene	3.9
123	Amy	3.9	654	Amy	3.9
456	Doris	3.9	123	Amy	3.9
456	Doris	3.9	876	Irene	3.9
456	Doris	3.9	654	Amy	3.9
567	Edward	2.9	765	Jay	2.9
...	...	...	...	...	...

How to exclude the same  
pairs in different order?

<> → < or >

# Table variables


---

```
SELECT S1.sID, S1.sName, S1.GPA, S2.sID, S2.sName, S2.GPA
FROM   Student S1, Student S2
WHERE  S1.GPA=S2.GPA AND S1.sID < S2.sID;
```

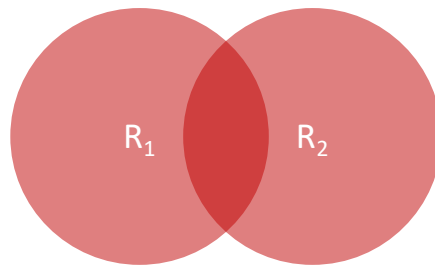
sID	sName	GPA	sID1	sName1	GPA1
123	Amy	3.9	456	Doris	3.9
123	Amy	3.9	876	Irene	3.9
123	Amy	3.9	654	Amy	3.9
456	Doris	3.9	876	Irene	3.9
456	Doris	3.9	654	Amy	3.9
567	Edward	2.9	765	Jay	2.9
654	Amy	3.9	876	Irene	3.9
543	Craig	3.4	789	Gary	3.4

# Set operators

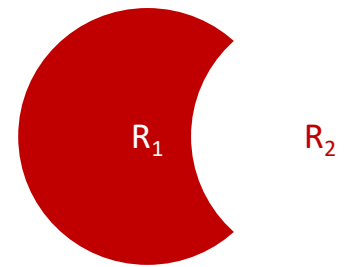
---



Union



Intersect



Except

# Sets, Bags and Lists

---

## Sets

Only one occurrence of each element

Unordered elements

## Bags (or multisets) $\rightarrow$ O que utilizamos

More than one occurrence of an element

Unordered elements and their occurrences

## Lists

More than one occurrence of an element

Occurrences are ordered

# Recall Bags

Multiset X

Tuple
(1, a)
(1, a)
(1, b)
(2, c)
(2, c)
(2, c)
(1, d)
(1, d)



$\lambda(X)$  = "Count of tuple in X"

Multiset X

Tuple	$\lambda(X)$
(1, a)	2
(1, b)	1
(2, c)	3
(1, d)	2

→ Número de  
ocorrências de  
cada elemento

In a set all counts are (0,1).



# Union as a bag operation

---

Multiset X

Tuple	$\lambda(X)$
(1, a)	2
(1, b)	0
(2, c)	3
(1, d)	0

U

Multiset Y

Tuple	$\lambda(Y)$
(1, a)	5
(1, b)	1
(2, c)	2
(1, d)	2

=

Multiset Z

Tuple	$\lambda(Z)$
(1, a)	7
(1, b)	1
(2, c)	5
(1, d)	2

$$\lambda(Z) = \lambda(X) + \lambda(Y)$$

Na união o número de ocorrências é a soma

# Intersect as a bag operation

---

Multiset X

Tuple	$\lambda(X)$
(1, a)	2
(1, b)	0
(2, c)	3
(1, d)	0

$\cap$

Multiset Y

Tuple	$\lambda(Y)$
(1, a)	5
(1, b)	1
(2, c)	2
(1, d)	2

$=$

Multiset Z

Tuple	$\lambda(Z)$
(1, a)	2
(1, b)	0
(2, c)	2
(1, d)	0

$$\lambda(Z) = \min(\lambda(X), \lambda(Y))$$

Na interseção o número de ocorrências é o mínimo

# Except as a bag operation

---

Multiset X

Tuple	$\lambda(X)$
(1, a)	2
(1, b)	0
(2, c)	3
(1, d)	0

—

Multiset Y

Tuple	$\lambda(Y)$
(1, a)	5
(1, b)	1
(2, c)	2
(1, d)	2

=

Multiset Z

Tuple	$\lambda(Z)$
(1, a)	0
(1, b)	0
(2, c)	1
(1, d)	0

If  $\lambda(X) > \lambda(Y)$

$$\lambda(Z) = \lambda(X) - \lambda(Y)$$

Else

0

College(cName, state, enr)  
Student(sID, sName, GPA, sizeHS)  
Apply(sID, cName, major, decision)

## Union operator

```
SELECT cName FROM College  
UNION  
SELECT sName FROM Student;
```

União

In SQL, the two sides of the union don't have to be the same

How to unify the two schemas?

```
SELECT cName AS name FROM College  
UNION  
SELECT sName AS name FROM Student;
```

Renomeamos para o mesmo

cName
Amy
Berkeley
Bob
Cornell
Craig
Doris
Edward
Fay
Gary
Helen
Irene
Jay
MIT
Stanford

# Union operator

College(cName, state, enr)

Student(sID, sName, GPA, sizeHS)

Apply(sID, cName, major, decision)

By default, in SQL, the union operator eliminates duplicates

If we want to have duplicates in our result

```
SELECT cName AS name FROM College
```

**UNION ALL → Devolve tudo independentemente de serem duplicados**

```
SELECT sName AS name FROM Student;
```

Result is not sorted anymore

Why? → O Algoritmo que utiliza para remover os duplicados implica ordenar as instâncias

How can we sort the result?

```
SELECT cName AS name FROM College
```

```
UNION ALL
```

```
SELECT sName AS name FROM Student
```

```
ORDER BY name;
```

name
Stanford
Berkeley
MIT
Cornell
Amy
Bob
Craig
Doris
Edward
Fay
Gary
Helen
Irene
Jay
Amy
Craig

# Kahoot time!

---

Any doubts?

# Readings

---

Jeffrey Ullman, Jennifer Widom, A first course in Database Systems 3<sup>rd</sup> Edition

Section 6.1 – Simple Queries in SQL

Section 6.2 – Queries Involving More Than One Relation

Section 6.3 - Subqueries

Section 6.4 – Full-Relation Operations

Section 6.5 – Database Modifications

Philip Greenspun, SQL for Web Nerds,  
<http://philip.greenspun.com/sql/>